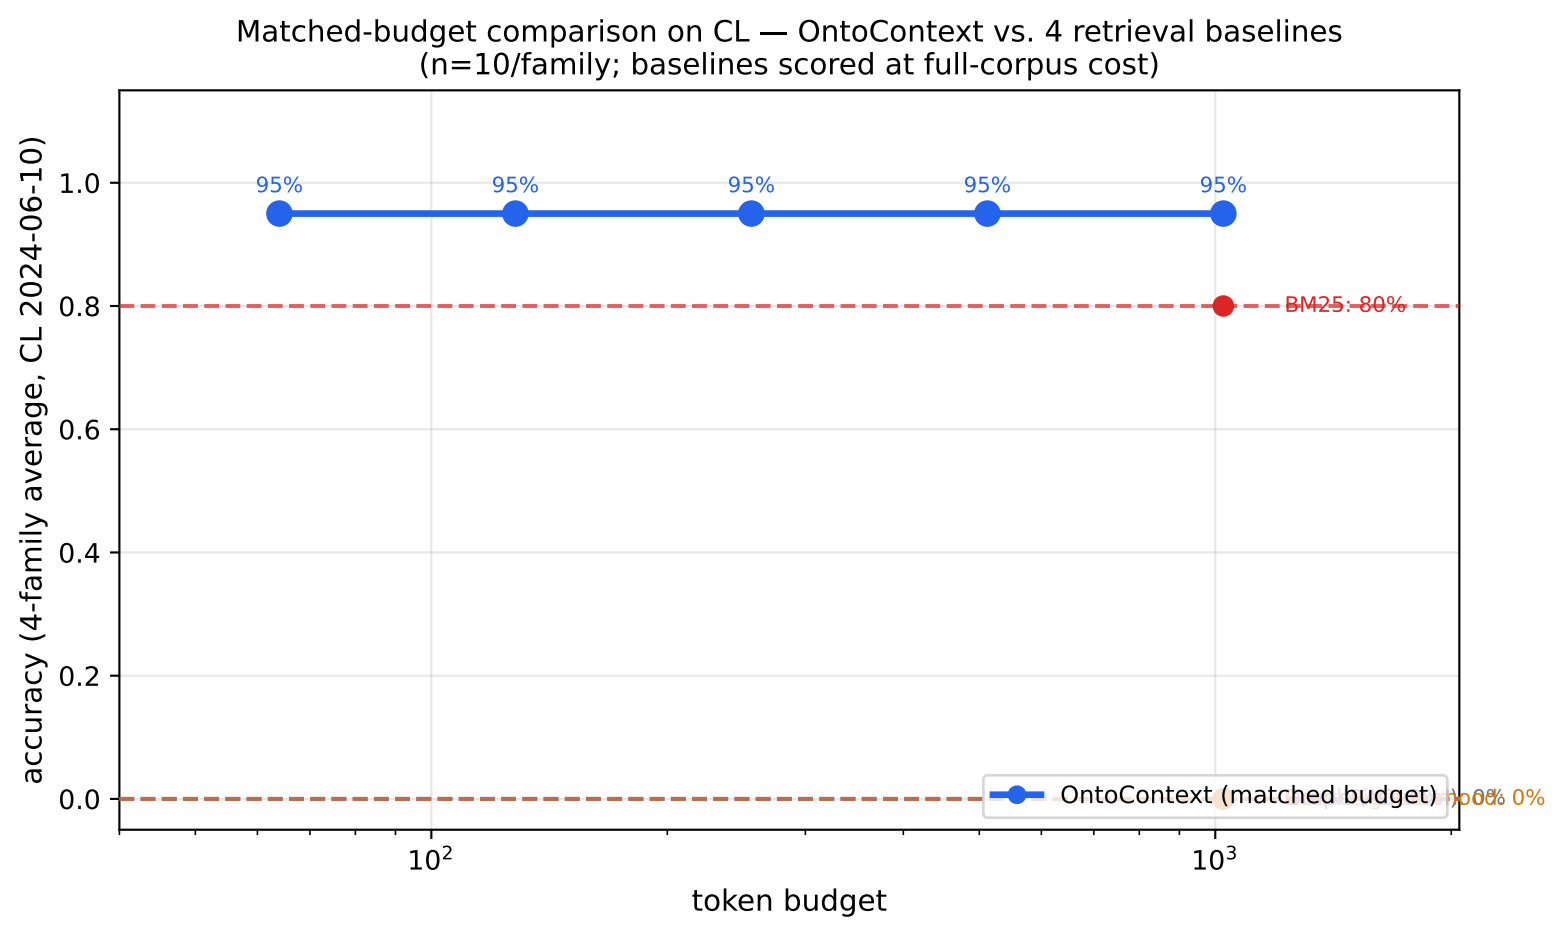
Matched-budget comparison on CL — OntoContext vs. 4 retrieval baselines
(n=10/family; baselines scored at full-corpus cost)
accuracy (4-family average, CL 2024-06-10)
1.0
0.8
0.6
0.4
0.2
0.0
95%
95%
95%
95%
95%
BM25: 80%
OntoContext (matched budget)
10^2
10^3
token budget